



INFS4205/7205 Advanced Techniques for High Dimensional Data

Managing High-Dimensional Data

Semester 1, 2021

University of Queensland

+ Advanced Techniques for High Dimensional Data

- Course Introduction
- Introduction to Spatial Databases
- Spatial Data Organization
- Spatial Query Processing
- Managing Spatiotemporal Data
- Managing High-Dimensional Data
- Introduction to Multimedia Database
- Route Planning in Road Network
- When AI Meets High-Dimensional Data
- Trends and Course Review

+ Outline

■ Motivation

- Examples
- Why
- What to expect

■ Technique

- X-Tree
- Pyramid
- VA-File
- iDistance
- Other techniques

■ Summary

+ Outline

■ Motivation

- Examples
- Why
- What to expect

■ Technique

- X-Tree
- Pyramid
- VA-File
- iDistance
- Other techniques


■ Summary

+ High-Dimensional Feature Space

5



- Picture elements in digital images
- Represented by $\langle \text{red, green, blue} \rangle$
 - Black: $\langle 0, 0, 0 \rangle$
 - Orange: $\langle 255, 165, 0 \rangle$



Pixel

$\text{Img} = \langle 7414, 230, 0, 0, \dots \rangle$



$2^8 = 256\text{D color space}$

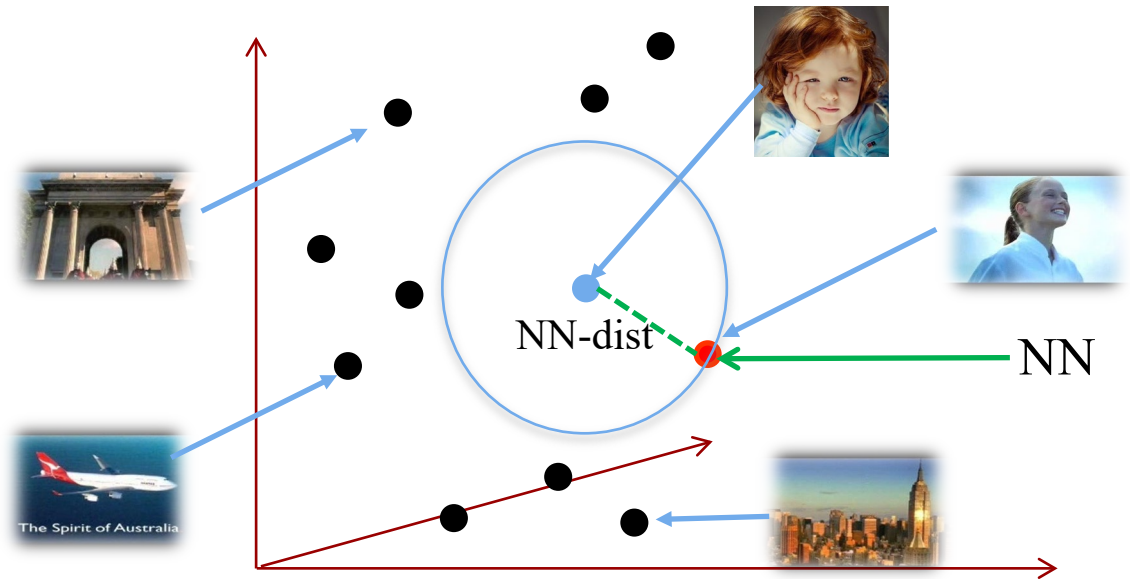
Red	Green	Blue	Pixel Count
0	0	0	7414
0	0	1	230
0	0	2	0
0	0	3	0
0	1	0	8
0	1	1	372
0	1	2	88
0	1	3	0
0	2	0	0
0	2	1	0
0	2	2	10

+ High-Dimensional Feature Space

6

$f_1 : \langle 0.336, 0.130, 0.023, 0.331, 0.132, 0.000, 0.120, 0.181 \rangle$
 $f_2 : \langle 0.331, 0.123, 0.028, 0.338, 0.008, 0.011, 0.132, 0.181 \rangle$
 $f_3 : \langle 0.331, 0.116, 0.028, 0.345, 0.101, 0.179, 0.133, 0.181 \rangle$
.....
 $f_n : \langle 0.331, 0.102, 0.021, 0.336, 0.009, 0.000, 0.009, 0.192 \rangle$

High-dimensional data points



High-D Feature Space

+ Outline

■ Motivation

- Examples
- Why
- What to expect

■ Technique

- X-Tree
- Pyramid
- VA-File
- iDistance
- Other techniques

■ Summary

+ CPU Cost for Indexing is High


- Indexing: to locate the data quickly
 - Query processing time should not increase linearly with the DB size
 - I/O cost: Disk page accesses
 - A balanced tree structure can reduce search time (I/O costs)
- However, for high-dimensional data
 - CPU cost: Computation of similarity/distance
 - The CPU cost for some basic operation is no longer negligible
 - As dimensionality increases, the portion of CPU cost in total response time increases
 - This is different from traditional databases, where only I/O costs are considered

+ Indexing vs. Linear Scan


- The performance of an index degrades rapidly as dimensionality increases, and eventually **underperforms linear scan!**
- However, linear scan needs to search the whole data file - affected volume is 100%
 - Then query processing time will increase at least linearly with the DB size
- When the size of dataset is very large, loading the whole data into memory is unlikely
 - Even if this is possible, it's still too expensive to scan all the data

+ Overlap in R* Tree

- Aims to minimize the overlap
 - Overlap: more than one branch need to be expanded
- Dimensionality and Overlap



a. Overlap (Uniformly Distributed Data)



b. Weighted Overlap (Real Data)



Figure 2: Overlap of R*-tree Directory Nodes depending on the Dimensionality

- (Overlap: % of space covered by more than one R*-tree node)
- (Weighted Overlap: % of data objects in overlapping space)

+ “Curse of Dimensionality”


- Too many partitions
- Too few points
- The nearest is not near enough

+ Too Many Partitions




- A simple partitioning scheme splits the data space in each dimension into 2 parts
 - Thus, d -dimensions implies 2^d partitions
- For a 1,000,000 (10^6) data points database:
 - If $d \leq 10$, at most 1024 partitions for 1M points
 - If d is large, e.g., $d = 100$, there are around 2^{100} (about 10^{30}) partitions for 10^6 data points (much more partitions than points!)
- An overwhelming majority of the partitions are empty!

+ Too Few Points (1)



Data space = $[0, 1]^d$



- Assume the data space is a hypercube, where each dimension is within the range of $[0, 1]$
- The probability of a uniformly distributed point p lying within a hypercube range query with side length s is
$$Pr[p \in \text{QueriedRegion}] = s^d$$
- For uniform data, when $d = 100$, a hypercube range query with $s = 0.95$ only covers 0.59% of the data points; compared to $d = 2$, 90.25%


+ Too Few Points (2)

- Most of the data lies close to the boundary of the dataspace
- Suppose we have a 50-dimensional hypercube with length of 1
 - Total Volume = $1 \times 1 \times 1 \times \dots \times 1 = 1^{50} = 1$
 - If we regard the inner 90% of each dimension as the interior region, and the outside of it as boundary region
 - $0.00 < x_1 < 0.90, 0.00 < x_2 < 0.90, \dots, 0.00 < x_{50} < 0.90$
 - If the data is uniformly distributed
 - The interior region's volume is $0.9^{50} \approx 0.005$
 - The boundary region's volume is $1 - 0.005 = 0.995$
 - 99.5% of the points are in the boundary region

+ Too Few Points (3)


15

- The probability of a uniformly distributed point lying within a spherical query $Sphere(q, 0.5)$ is:



- $d = 2$: $\pi r^2 / (2r)^2 = \pi/4 = 0.785$
- $d = 3$: $(\frac{4\pi}{3} \times r^3) / (2r)^3 = \pi/6 = 0.524$
- ...
- $d = 8$: $(\frac{\pi^4}{24} \times r^8) / (2r)^8 = \pi^4 / (3 \times 2^{11}) = 0.016$

- Note that $Sphere(q, 0.5)$ is the largest spherical query that fits within the unit data space




+ The Nearest is not Near Enough

- In some sense, nearly all of the uniformly distributed high-dimensional points are **"far away"** from the centre
 - The high-dimensional data are almost entirely in the "corners" of the hypercube, with almost nothing in the "middle".
- **Low distance contrast**: as the dimensionality increases, the distance to the nearest neighbour **approaches** the distance to the furthest neighbour.
 - For uniformly distributed data, difference between the nearest data point and the furthest data point reduces greatly


+ The Nearest is not Near Enough (2)

17

■ Expected NN distance vs.



dimensionality



number of data points

+ Question:

- Is the nearest neighbour meaningful in a high-D space?
- It depends on the data distribution
 - Nothing can be done for uniformly distributed data.
 - Generally, real data are not uniformly distributed, but exhibit certain clusters, trends or skewness.

+ Outline

■ Motivation

- Examples
- Why
- What to expect

■ Technique

- X-Tree
- Pyramid
- VA-File
- iDistance
- Other techniques

■ Summary

+ Design Principles

■ Simple

- Simple in design
- Easy to be integrated into existing DBMS
 - It's very hard to add a new indexing structure into a DBMS
 - Better chance to be practically useful if building on top of a mature and commercially available indexing method (e.g.. B+-tree & R-tree)

■ Efficient

- Efficient in disk access (I/O cost) /CPU time
- Must support efficient updates (insert/delete/update operations)

+ Outline

■ Motivation

- Examples
- Why
- What to expect

■ Technique

- X-Tree
- Pyramid
- VA-File
- iDistance
- Other techniques

■ Summary

+ Outline

■ Motivation

- Examples
- Why
- What to expect

■ Technique

- X-Tree
- Pyramid
- VA-File
- iDistance
- Other techniques


■ Summary

+ X-Tree

- Variation of R-Tree and compromise between the hierarchy and sequential access
- Avoid overlap in the directory, when we cannot find a good partition
 - In high dimensional, linear organization of the directory sometimes is more efficient
 - Due to the high overlap, most of the directory has to be searched
 - So don't split when its overlap is high

■ Tree Nodes

- Data Node
- Normal Directory Node
- Super Node
 - Large directory nodes of variable size
 - Arbitrary number of blocks
 - Avoid split in the directory that would result in an inefficient directory



+ X-Tree

■ Insertion

- If no split occurs, update the size of MBR
- If split is required
 1. If the overlap is low
 - Split the node with same techniques as R-tree Based techniques (or other techniques)
 2. Otherwise (the overlap is high), do overlap minimal split
 - If the number of MBRs in one of the partitions is below a given threshold (unbalanced)
 - Stop splitting and extend to a Supernode

+ Overlap Minimal Split*

- Split History

- The dimension according to which an MBR has been split
- Which new MBRs have been created by this split


- Overlap Minimal Split

- For point data
 - It is always possible to find an **overlap-free** split
 - It is not possible to guarantee that the two sets are balanced
 - Determine a dimension according to which all MBRs have been split previously

+ Overlap Minimal Split for Point Data*

■ Split Tree

- Leaf Node: MBR
- Internal Node: Nodes that have been split into new MBRs
 - Record the dimension that was used to split



- All MBRs in any split tree have one split dimension in common: the root node

+ Outline

■ Motivation

- Examples
- Why
- What to expect



■ Technique

- X-Tree
- Pyramid
- VA-File
- iDistance
- Other techniques

■ Summary

+ Pyramid Technique

- Divide the data space such that the resulting partitions are shaped like peels of an onion
- Divide the d -D space into $2d$ pyramids having the center point $(0.5, 0.5, \dots, 0.5)$ of the space as their top
- Each pyramid is cut into slices parallel to the basis
- Data is mapped from d -dimensional space to 1-dimensional space
- B+-tree can then be used
- Such that typical window query will not overlap all the nodes as the balanced split




+ Pyramid Numbering

■ Pyramid Numbering: a point v belongs to pyramid i

■ Determine i

■ j_{max} : The dimensional j that has the maximum deviation $|0.5 - v_j|$

$$i = \begin{cases} j_{max} & , \text{if } v_{j_{max}} < 0.5 \\ j_{max} + d, & \text{if } v_{j_{max}} \geq 0.5 \end{cases}$$






+ Point Height in a Pyramid

30


■ Height of v

■ The distance from the point to the center according to dimension j_{max}

$$h_v = |0.5 - v_i \bmod d|$$




+ Pyramid Value




■ Pyramid Value of d -dimensional point v

- Transform a d -dimensional point to a value $pv_v = i + h_v$
 - i is an integer
 - h_v is in range $[0, 0.5]$
 - pv_v is in range $[i, i + 0.5]$
- Values covered by different pyramids are disjoint



$$pv_v = 1 + 0.3 = 1.3$$




- Use B⁺-Tree to index
 - Easy to insert, delete, and update

+ Pyramid Technique

■ Query is complex

- Point Query: Compute the pyramid value and query the B⁺-Tree
- Range Query $[q_{0_{min}}, q_{0_{max}}], \dots, [q_{d-1_{min}}, q_{d-1_{max}}]$
 1. Determine the affected pyramids
 - Transform one d -dimensional range query q into an equivalent $2d$ range queries, one for each pyramid
 2. Determine the ranges inside the pyramids
 - $[i + h_{low}, i + h_{high}]$ for each pyramid




+ Pyramid Technique


33

- The effectiveness is sensitive to query position

- Query region vs. search region



- Non-uniform distribution – design non-uniform pyramid



+ Outline

■ Motivation

- Examples
- Why
- What to expect

■ Technique

- X-Tree
- Pyramid
- VA-File
- iDistance
- Other techniques

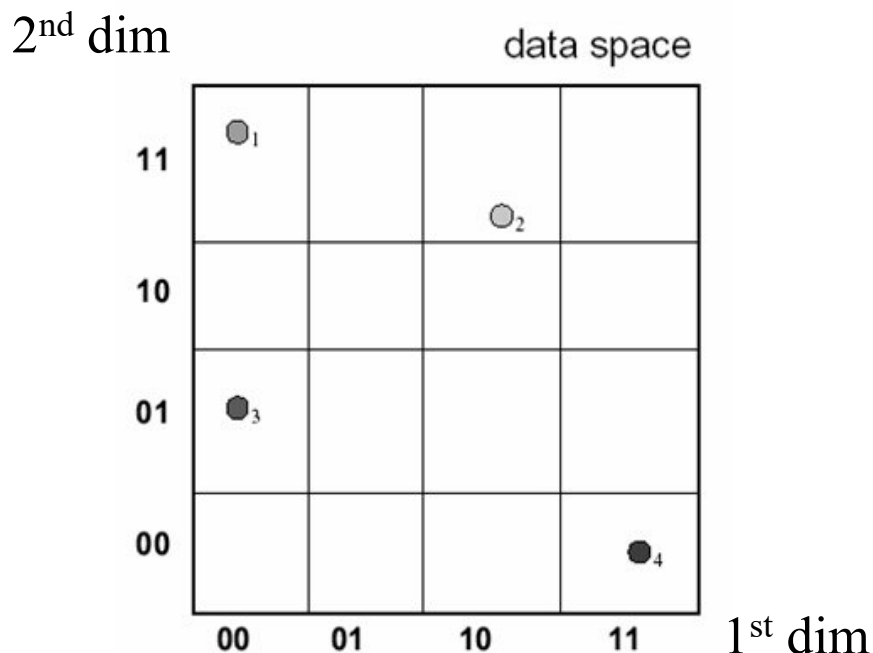
■ Summary

+ Vector Approximation (VA)-File

- In High-D spaces, tree-based indexing structures may examine a very large fraction of leaf nodes
 - Since the MBRs heavily overlap if data are distributed uniformly
- Let go of the hierarchies and sequentially scan the whole data set
 - Sequential scan is much faster than random read due to the cost of disk seek operations
- Natural question: how to speed up linear scan?
 - Using approximation to compress vector data
 - Easing the computation and reducing the amount of data to exam during the search

+ Basic Idea

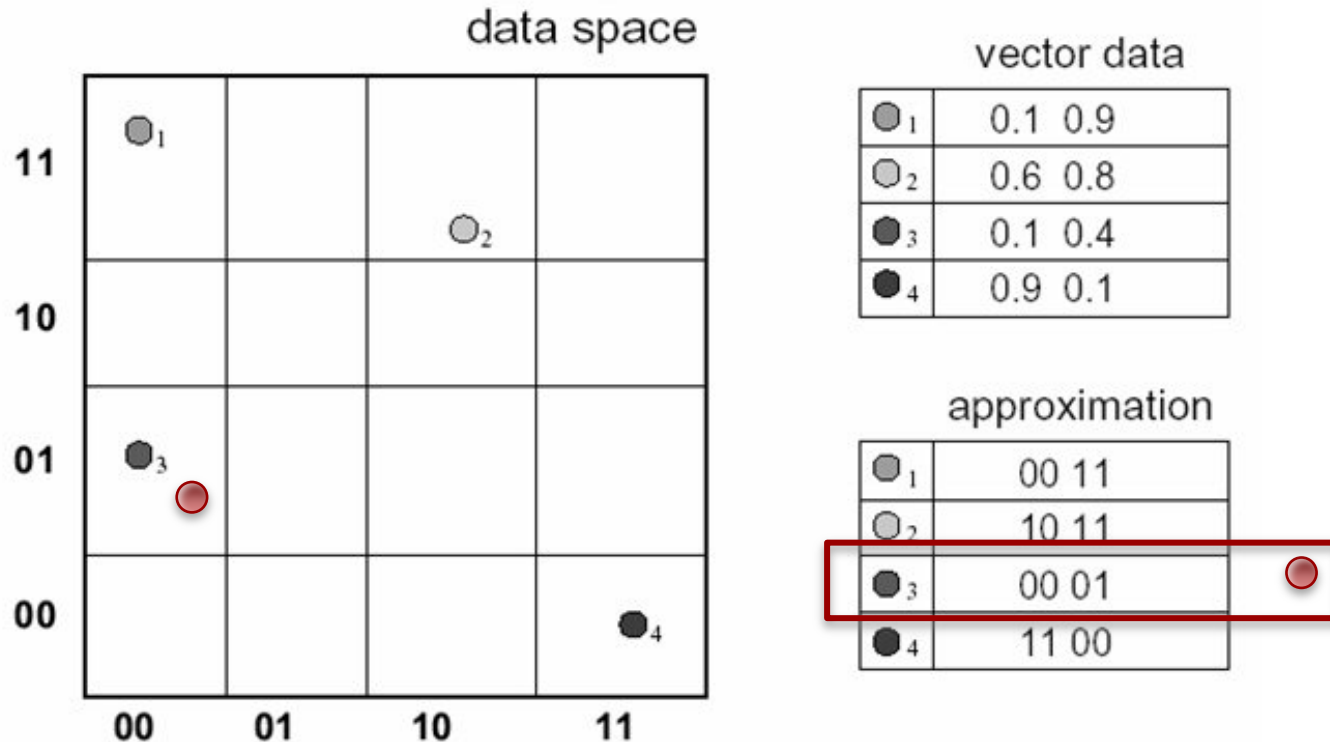
- For each dimension i , a small number of bits b_i is assigned to divide the dimension into 2^{b_i} intervals
- Let b be the sum of all dimensions' b_i 's,
 - The data space is divided into 2^b cells
 - Every point is represented by b bits



- ✓ 1st dim: 2 bits
- ✓ 2nd dim: 2 bits
- ✓ In total $2^2 \times 2^2 = 2^{2+2} = 2^4 = 16$ cells

+ Building VA-file and Query

■ Map objects to cells

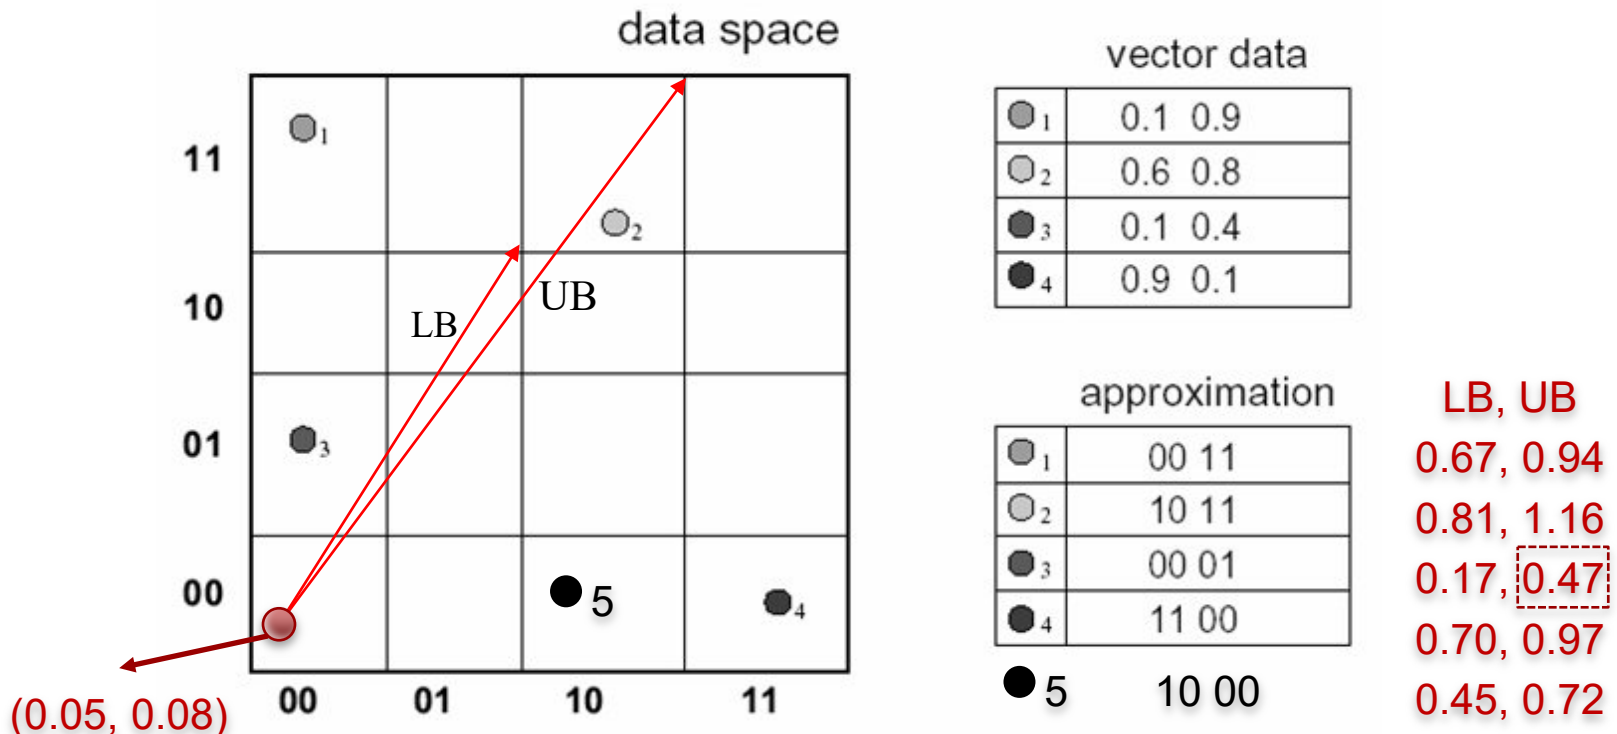


- Each cell has a bit representation with length b
- The VA-file itself is simply an array of bits concatenation based on the quantization of the original feature vectors.

+ VA-file 2-Phase NN search

■ Phase 1: Filtering


- The VA-file is **sequentially** scanned.
- The lower and upper bounds on the distance for each object's approximation (i.e., VA) is then computed.
- Assume δ is the smallest upper bound so far, eliminate approximations with a lower bound that exceeds δ



+ VA-file 2-Phase NN search

■ Phase 1: Filtering

- The VA-file is **sequentially** scanned.
- The lower and upper bounds on the distance for each object's approximation (i.e., VA) is then computed.
- Assume δ is the smallest upper bound so far, eliminate approximations with a lower bound that exceeds δ

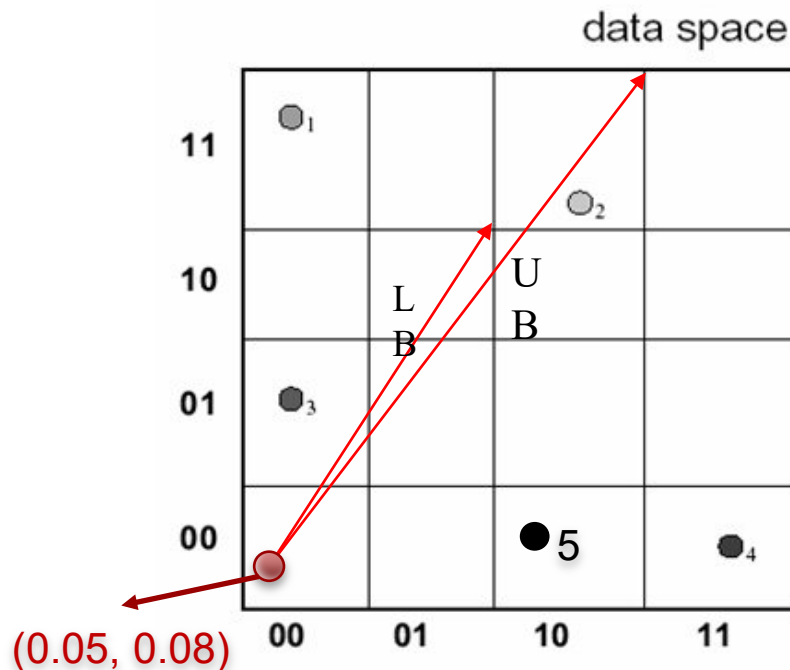


+ VA-file 2-Phase NN search

40

■ Phase 2: Refine

- After the filtering step, a small set of candidates remain
- Candidates are sorted in ascend order by lower bound P3, P5
- Calculate the current NN distance by iteration through the sorted candidates
- If a lower bound is encountered that exceeds the nearest distance seen so far, the VA-file method stops



vector data

● ₁	0.1 0.9
● ₂	0.6 0.8
● ₃	0.1 0.4
● ₄	0.9 0.1

approximation

● ₁	00 11
● ₂	10 11
● ₃	00 01
● ₄	11 00
● ₅	10 00

LB, UB	Dis.
0.67, 0.94	
0.81, 1.16	
0.17, 0.47	0.32
0.70, 0.97	
0.45, 0.72	

+ What We can Learn from VA File?

- How to speed up linear scan ?

Answer: Use approximation!

- Use only b_i bits per dimension
 - Floating points: 32bits per dimension
 - Speed up the scan by a factor of $32/b_i$
- Identify all points which could be returned as an answer
- Verify the points by accessing the original feature vectors

+ Limits of VA-file

- The total cost of VA-file also includes the **random access (I/O)** to the candidates. If the candidate set is large, save in linear scan on VA-file will be offset.
- Large b_i :
 - Small number of candidates, but
 - Large VA-file
- VA-file performs best in **uniformly** distribution data. However, real data exhibit certain degree of skewness.
 - Assume the features are independent
 - Because the slices are obtained independently
 - VA-file divides the dimensions either for equal size or for equal population.

+ Outline

■ Motivation

- Examples
- Why
- What to expect

■ Technique

- X-Tree
- Pyramid
- VA-File
- iDistance
- Other techniques

■ Summary

+ iDistance^[1]

- iDistance is simple, but efficient
- It is a **distance and partition** based index
- It can be used for approximate search
- The index can be integrated to existing systems easily
- Another representative way to handle high dimensional data

+ Basic Ideas

■ Observations

1. The similarity/dissimilarity between points can be derived with reference to a chose **reference point**
2. Points can be **ordered** based on their distance to the reference point
3. Distance is essentially a single dimensional value
 - Reuse the existing 1-D index like B+ -Tree


+ Basic Ideas

- Data points are partitioned into clusters/partitions
- Points are transformed into 1-D space
 1. The high-dimensional space is split into partitions
 - P_0, P_1, \dots, P_{m-1}
 2. A reference point is identified for each partition
 - O_0, O_1, \dots, O_{m-1}
 3. A data point $p(x_1, x_2, \dots, x_d)$ in the i^{th} partition can be referenced via O_i in terms of the distance from p to O_i
 - Index key y : $y = i \times c + dist(p, O_i)$
 - c : a constant of the data range
 - All points in partition P_i mapped to range $[i \times c, (i + 1) \times c]$
- Data points are indexed based on similarity (metric distance) to such a point using a standard B⁺-tree

+ Indexing Points Based on Similarity

47

Indexing points based on similarity



Assume: $\text{Dist}(p_1, O_1) = 5$, $\text{Dist}(p_4, O_2) = 5$

We set $c = 100$

$i\text{Dist}(p_1) = 105$ and $i\text{Dist}(p_4) = 205$

+ The range of search NN

48

- The triangle inequality

$$\text{dist}(O_i, q) - \text{dist}(p, q) \leq \text{dist}(O_i, p) \leq \text{dist}(O_i, q) + \text{dist}(p, q)$$

- Consider the points within a range around q containing p


$$\text{dist}(O_i, q) - \text{querydist}(q) \leq \text{dist}(O_i, p) \leq \text{dist}(O_i, q) + \text{querydist}(q)$$

- The maximum distance range

$$\text{dist}(O_i, p) \leq \text{dist_max}_i$$

The final search space are points within the range

$$[\text{dist}(O_i, q) - \text{querydist}(q), \min(\text{dist_max}_i, \text{dist}(O_i, q) + \text{querydist}(q))]$$



+ The range of search NN

- The triangle inequality

$$\text{dist}(O_i, q) - \text{dist}(p, q) \leq \text{dist}(O_i, p) \leq \text{dist}(O_i, q) + \text{dist}(p, q)$$

- Consider the points within a range around q containing p


$$\text{dist}(O_i, q) - \text{querydist}(q) \leq \text{dist}(O_i, p) \leq \text{dist}(O_i, q) + \text{querydist}(q)$$

- The maximum distance range

$$\text{dist}(O_i, p) \leq \text{dist_max}_i$$

The final search space are
points within the rage

$$[\text{dist}(O_i, q) - \text{querydist}(q), \\ \min(\text{dist_max}_i, \text{dist}(O_i, q) + \text{querydist}(q))]$$



+ The range of search NN

50

- The triangle inequality

$$\text{dist}(O_i, q) - \text{dist}(p, q) \leq \text{dist}(O_i, p) \leq \text{dist}(O_i, q) + \text{dist}(p, q)$$

- Consider the points within a range around q containing p


$$\text{dist}(O_i, q) - \text{querydist}(q) \leq \text{dist}(O_i, p) \leq \text{dist}(O_i, q) + \text{querydist}(q)$$

- The maximum distance range

$$\text{dist}(O_i, p) \leq \text{dist_max}_i$$


The final search space are
points within the rage

$$[\text{dist}(O_i, q) - \text{querydist}(q), \\ \min(\text{dist_max}_i, \text{dist}(O_i, q) + \text{querydist}(q))]$$



+ Example

51



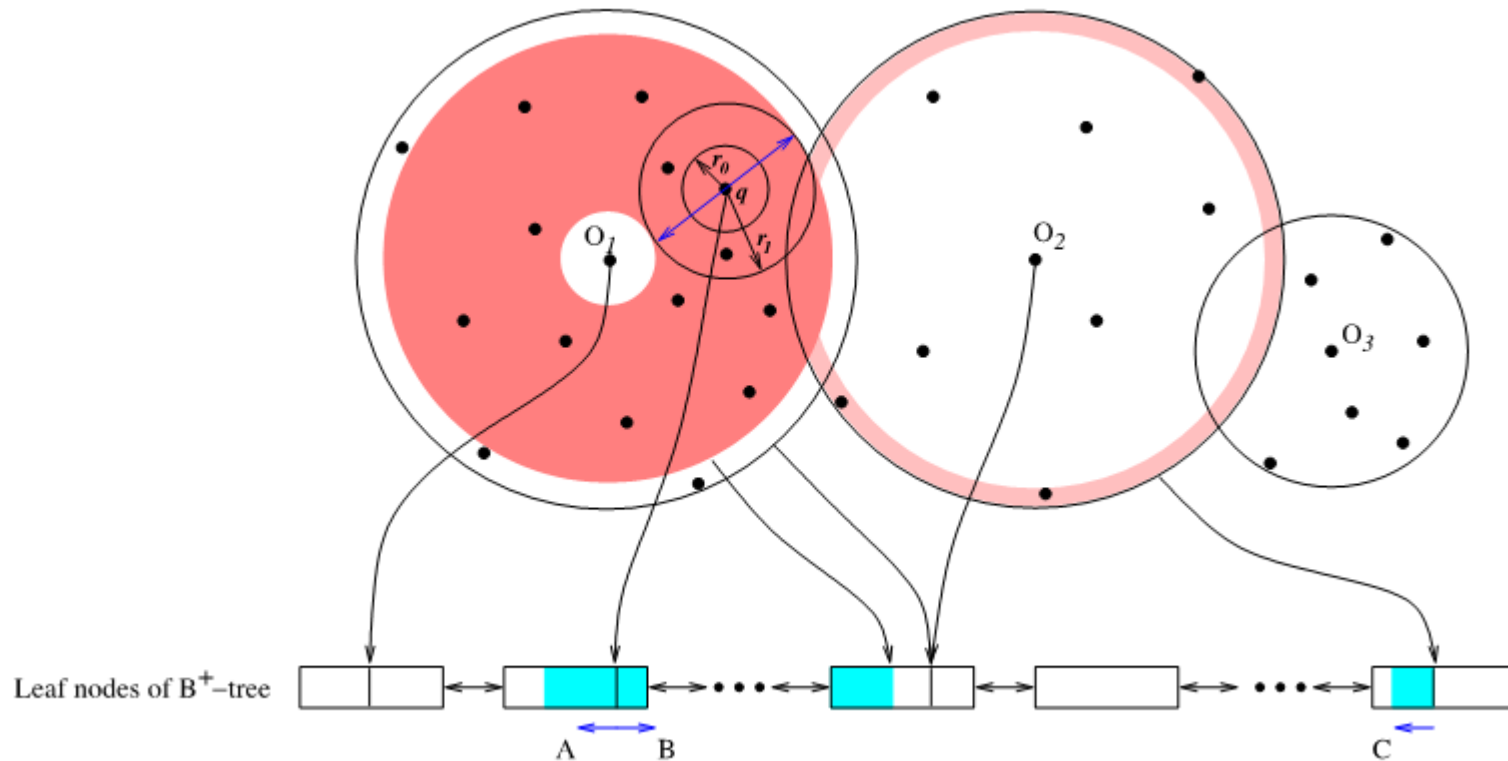
- Notice: all points along the same radius have the same value after transformation to distance
- Points with the same values are not necessarily close to each other

+ kNN Search

- Query distance $r = querydist$ is enlarged incrementally
- For each query distance r , search is conducted around the query point q and overlapped partition until kNN is found
- For each partition, its **minimal** and **maximal** distance to the Reference Point are recorded in the auxiliary structure
 - The auxiliary structure: $\{(O, \min, \max)\}$
- Scanning the auxiliary structure to identify the partitions whose data space overlaps with the query sphere (q, r)

+ Graphical Illustration of kNN search

53



A range in B+-tree

- ✓ Searching region is enlarged until getting kNN

+ Formation of Clusters/Partitions

■ Equal Partitioning:

- Effective if data is uniformly distributed
- However, data points are often non-uniformly distributed

■ Number of Clusters:

- Small number: More points are likely to have similar distance to a given reference point
- Large number: More circles are likely to overlap and incurs additional traversal and searching

+ Outline

■ Motivation

- Examples
- Why
- What to expect

■ Technique

- VA-File
- X-Tree
- Pyramid
- iDistance
- Other techniques...

■ Summary

+ List of techniques

- Hierarchical Tweaking
 - **X-Tree** [VLDB'96]
 - SS-Tree, SR-Tree, Ball-Tree
- Transformation based
 - **Pyramid Technique** [SIGMOD'98]
 - **iDistance** [VLDB'01]
 - Optimal one-dimensional distance B+ tree [SIGMOD'05]
- Dimensionality reduction
 - Local Dimensionality Reduction (LDR)
 - Multi-level Mahalanobis based Dimensionality Reduction (MMDR)
- Data compression based
 - Vector Approximation (**VA-File**) [VLDB'98]
- Hybrid of tree-like structure and data compression
 - Independent Quantization (IQ-tree)
 - Local Digital Coding (LDC)
- Approximate search
 - Locality Sensitive Hashing (LSH)
 - Vector Quantization (VQ-index)
 - Spatial Approximation Sample Hierarchy (SASH)

+ Outline

■ Motivation

- Examples
- Why
- What to expect

■ Technique

- VA-File
- X-Tree
- Pyramid
- iDistance
- Other techniques...

■ Summary

+ Summary

- “Curse of Dimensionality” is a real problem which is very difficult to solve
- It is counter-intuitive that similarity search can be meaningless with more information captured
- It is also counter-intuitive that sophisticated indexing structures may not even be as efficient as a linear scan over the entire database
- People often use pivot-points based indexing and data compression to deal with high dimensional data
- This is still an open problem, with most solutions which are either data dependent or domain specific

+ Readings

- Christian Böhm, Stefan Berchtold and Daniel Keim: “Searching in High-dimensional Spaces: Index Structures for Improving the Performance of Multimedia Databases”, *ACM Computing Surveys* 33 (3), 2001.
- Other papers on X-Tree, Pyramid Technique, VA-file and iDistance.

+ Advanced Techniques for High Dimensional Data

60

- Course Introduction
- Introduction to Spatial Databases
- Spatial Data Organization
- Spatial Query Processing
- Managing Spatiotemporal Data
- Managing High-Dimensional Data
- Introduction to Multimedia Database-next week
- Route Planning in Road Network
- When AI Meets High-Dimensional Data
- Trends and Course Review