

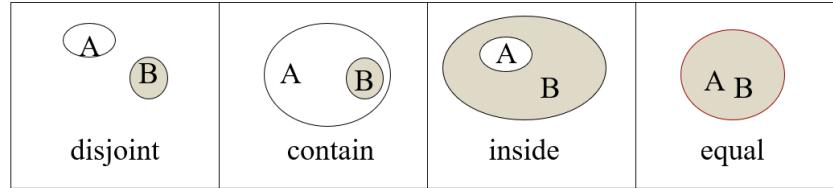
Tutorial 1: Spatial Database Introduction

Semester 1, 2021

Question 1: The 9-Intersection Matrix is a topological model and a standard used to describe the spatial relations of two polygons (although it can be extended to line, we only discuss polygon here). Each polygon has three parts: A boundary, an interior, and an exterior. Therefore, for any two polygons, their topological relationships can be characterized using nine possible intersections between their interiors/boundaries/exteriors.

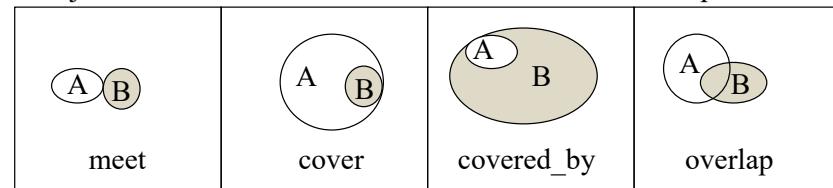
$$\begin{pmatrix} A_b \cap B_b & A_b \cap B_i & A_b \cap B_e \\ A_i \cap B_b & A_i \cap B_i & A_i \cap B_e \\ A_e \cap B_b & A_e \cap B_i & A_e \cap B_e \end{pmatrix}$$

Depending on the *True/False* of each intersection, eight and only eight topological relations can be identified. Please discuss and understand them.



$$\begin{pmatrix} 0 & 0 & 1 \\ 0 & 0 & 1 \\ 1 & 1 & 1 \end{pmatrix} \quad \begin{pmatrix} 0 & 0 & 1 \\ 1 & 1 & 1 \\ 0 & 0 & 1 \end{pmatrix} \quad \begin{pmatrix} 0 & 1 & 0 \\ 0 & 1 & 0 \\ 1 & 1 & 1 \end{pmatrix} \quad \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

disjoint contain inside equal



$$\begin{pmatrix} 1 & 0 & 1 \\ 0 & 0 & 1 \\ 1 & 1 & 1 \end{pmatrix} \quad \begin{pmatrix} 1 & 0 & 1 \\ 1 & 1 & 1 \\ 0 & 0 & 1 \end{pmatrix} \quad \begin{pmatrix} 1 & 1 & 0 \\ 0 & 1 & 0 \\ 1 & 1 & 1 \end{pmatrix} \quad \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix}$$

meet cover covered_by overlap

Question 2: A demonstration of spatial database using PostGIS

PostGIS is a spatial extension that adds support for geographic objects to the PostgreSQL (Just like Oracle Spatial to Oracle). This tutorial works as a demonstration of some basic spatial operations. Some of the useful tools and link:

- PostGIS, you can download and try it on your computer : <https://postgis.net>
- The complete tutorial can here: <https://postgis.net/workshops/postgis-intro/>
- To view the spatial data, you can use the free QGIS: <http://qgis.org>
- The New York dataset used in this tutorial: <http://s3.cleverelephant.ca/postgis-workshop-2018.zip>

Spatial Database Creation

After the installation, open the pgsql command line to interact with PostgreSQL

- Create a database
CREATE DATABASE sdbms;
- Connect to the new database
 \c sdbms
- Load the PostGIS spatial extension
 Create EXTENSION postgis;
- Confirm that PostGIS is installed
 SELECT postgis_full_version();

```
jinfm=# CREATE DATABASE sdbms;
CREATE DATABASE
jinfm=# \c sdbms
You are now connected to database "sdbms" as user "jinfm".
sdbms=# CREATE EXTENSION postgis;
ERROR:  extension "postgis" already exists
sdbms=# SELECT postgis_full_version();

postgis_full_version
-----
-----
```



```
-----  
-----  
-----  
POSTGIS="3.1.1 aaf4c79" [EXTENSION] PGSQL="130" GEOS="3.8.1-CAPI-1.13.3" PROJ="7.1.1"  
GDAL="GDAL 3.1.4, released 2020/10/20" LIBXML="2.9.10" LIBJSON="0.13.1" LIBPROTOBUF="1.  
3.3" WAGYU="0.5.0 (Internal)" TOPOLOGY RASTER  
(1 row)
```

Data Import

The data used in this tutorial is the format of “shapefile”, which commonly refers to a collection of files with .shp, .shx, .dbf that have the same prefix name. A shapefile is a simple, nontopological format for storing the geometric location and attribute information of geographic features.

Geographic features in a shapefile can be represented by points, lines, or polygons (areas). The workspace containing shapefiles may also contain dBASE tables, which can store additional attributes that can be joined to a shapefile's features.

```

POINT(0 0)
MULTIPOINT((0 0),(1 2))

LINESTRING(0 0,1 1,1 2)
MULTILINESTRING((0 0,1 1,1 2),(2 3,3 2,5 4))

```

```

POLYGON((0 0,4 0,4 4,0 4,0 0),(1 1,2 1,2 2,1 2,1 1))
MULTIPOLYGON(((0 0,4 0,4 4,0 4,0 0),(1 1,2 1,2 2,1 2,1 1)),((-1 -1,-1 -2,-2 -2,-2 -1,-1 -1)))

```

For example, the three mandatory files of the nyc_street shapefile are:

- **nyc_street.shp**: The main shape file, store the feature geometry like points, lines, and polygons (areas). The area features are represented as closed loop, double-digitized polygons. All the geometries are stored as a set of vector coordinates.
- **nyc_street.shx**: The shape index file, stores the offset of the corresponding main file record from the beginning of the main file.
- **nyc_street.dbf**: The attribute table file, stores the columnar attributes for each shape in dBASE format. It can be joined to a shapefile's features.

More details of the shapefile can be found here:

http://downloads.esri.com/support/whitepapers/mo_shapefile.pdf

To import a shapefile into the database, we first use shp2pgsql command line tool that shipped with PostGIS to convert it into SQL script. For example:

- shp2pgsql -s 26918 nyc_streets.shp nyc_streets > nyc_streets.sql

The -s parameter set the *Spatial Reference Identifier* (SRID), which defines all the parameters of the data's geographic coordinate system and projection. All the other shapefiles can be converted similarly.

```
[jinfm@d-i184-53-246 ~ % cd Desktop/INFS_4205_7205_tutorial/week1/postgis-workshop-2018/data
[jinfm@d-i184-53-246 data % shp2pgsql -s 26918 nyc_streets.shp nyc_streets > nyc_streets.sql
Field id is an FTDouble with width 11 and precision 0
Shapefile type: Arc
Postgis type: MULTILINESTRING[2]
```

After that, in our sdmbs, we import these SQL scripts using command \i

- \i path\to\nyc_streets.sql

```
[sdbms=# \i /Users/jinfm/Desktop/INFS_4205_7205_tutorial/week1/postgis-workshop-2018/data/nyc_streets.sql
SET
SET
BEGIN
CREATE TABLE
ALTER TABLE
    addgeometrycolumn
-----
public.nyc_streets.geom SRID:26918 TYPE:MULTILINESTRING DIMS:2
(1 row)

INSERT 0 1
```

- \d nyc_streets

```
[sdbms=# \d nyc_streets
Table "public.nyc_streets"
 Column |          Type          | Collation | Nullable | Default
-----+----------------+-----+-----+-----+
 gid   | integer          |          | not null | nextval('nyc_streets_gid_seq'::regclass)
 id    | double precision   |          |          | 
 name  | character varying(200)|          |          | 
 oneway| character varying(10)|          |          | 
 type  | character varying(50)|          |          | 
 geom  | geometry(MultiLineString,26918) |          |          | 
```

Indexes:

```
"nyc_streets_pkey" PRIMARY KEY, btree (gid)
```

Now you can use QGIS to connect to your database and see these geometries.



Data Description

Before we query into the data, let's first describe their attributes:

nyc_census_blocks

A census block is the smallest geography for which census data is reported. All higher level census geographies (block groups, tracts, metro areas, counties, etc) can be built from unions of census blocks.

blkid	A 15-digit code that uniquely identifies every census block . Eg: 360050001009000
popn_total	Total number of people in the census block
popn_white	Number of people self-identifying as "White" in the block
popn_black	Number of people self-identifying as "Black" in the block
popn_nativ	Number of people self-identifying as "Native American" in the block
popn_asian	Number of people self-identifying as "Asian" in the block
popn_other	Number of people self-identifying with other categories in the block
boroname	Name of the New York borough. Manhattan, The Bronx, Brooklyn, Staten Island, Queens
geom	Polygon boundary of the block

nyc_neighborhoods

name	Name of the neighborhood
boroname	Name of the New York borough. Manhattan, The Bronx, Brooklyn, Staten Island, Queens
geom	Polygon boundary of the neighborhood

nyc_streets

name	Name of the street
oneway	Is the street one-way? "yes" = yes, "" = no
type	Road type (primary, secondary, residential, motorway)
geom	Linear centerline of the street

nyc_subway_stations

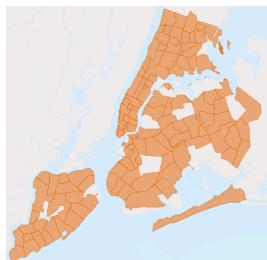
name	Name of the station
borough	Name of the New York borough. Manhattan, The Bronx, Brooklyn, Staten Island, Queens
routes	Subway lines that run through this station
transfers	Lines you can transfer to via this station
express	Stations where express trains stop, “express” = yes, “” = no
geom	Point location of the station

```
[sdbms=# \dt
      List of relations
 Schema |           Name            | Type | Owner
-----+-----+-----+-----+
 public | nyc_census_blocks    | table | jinfm
 public | nyc_neighborhoods   | table | jinfm
 public | nyc_streets          | table | jinfm
 public | nyc_subway_stations | table | jinfm
 public | spatial_ref_sys     | table | jinfm
(5 rows)
```

QGIS for visualization



nyc_census_blocks



nyc_neighborhoods



nyc_streets



nyc_subway_station

Simple Spatial SQL

1. What is the population of the City of New York?

```
SELECT Sum(popn_total) AS population
FROM  nyc_census_blocks;
```

```
[sdbms=# Select Sum(popn_total) As population
[sdbms-# From nyc_census_blocks;
 population
-----
 8175032
(1 row)
```

2. What is the population of the Bronx?

```
SELECT Sum(popn_total) AS population  
FROM nyc_census_blocks  
WHERE boroname = 'The Bronx'
```

```
[sdbms=# Select Sum(popn_total) As population  
[sdbms=# From nyc_census_blocks  
[sdbms=# Where boroname='The Bronx';  
population  
-----  
1385108  
(1 row)
```

3. For each borough, what percentage of the population is white?

```
SELECT boroname, 100 * Sum(popn_white)/Sum(popn_total) AS white_pct  
FROM nyc_census_blocks  
GROUP BY boroname;
```

```
[sdbms=# Select boroname, 100*Sum(popn_white)/Sum(popn_total) As white_pct  
[sdbms=# From nyc_census_blocks  
[sdbms=# Group By boroname;  
boroname | white_pct  
-----+-----  
Queens | 39.72207739459101  
Brooklyn | 42.80117379326865  
The Bronx | 27.903744689944755  
Manhattan | 57.44930394804628  
Staten Island | 72.8942034860154  
(5 rows)
```

Geometry Query

1. What is the area of the ‘West Village’ neighborhood?

```
SELECT ST_Area(geom) ST_Area(): Returns the area of a polygonal geometry.  
FROM nyc_neighborhoods  
WHERE name = 'West Village';
```

```
[sdbms=# Select ST_Area(geom)  
[sdbms=# From nyc_neighborhoods  
[sdbms=# Where name='West Village';  
st_area  
-----  
1044614.5296485956  
(1 row)
```

2. What is the area of Manhattan in acres?

```
SELECT Sum(ST_Area(geom)) / 4047  
FROM nyc_neighborhoods  
WHERE boroname = 'Manhattan';
```

```
[sdbms=# Select Sum(ST_Area(geom))/4047
[sdbms=# From nyc_neighborhoods
[sdbms=# Where boroname='Manhattan';
    ?column?
-----
13965.32012239119
(1 row)
```

ST_GeometryN(geom,n):
Returns the n-th geometry within a geometry object.

3. How many census blocks in New York City have a hole in them?

```
SELECT Count(*)
FROM nyc_census_blocks
WHERE ST_NumInteriorRings(ST_GeometryN(geom,1))>0;
```

```
[sdbms=# Select Count(*)
[sdbms=# From nyc_census_blocks
[sdbms=# Where ST_NumInteriorRings(ST_GeometryN(geom,1))>0;
    count
```

ST_NumInteriorRings(geom):

```
-----
```

```
43
(1 row)
```

Return the number of interior rings of the first polygon in the geometry.

This will work with both POLYGON and MULTIPOLYGON types but only looks at the first polygon.
Return NULL if there is no polygon in the geometry.

4. What is the total length of streets (in kilometers) in New York City?

```
SELECT Sum(ST_Length(geom)) / 1000
FROM nyc_streets;
```

```
[sdbms=# Select Sum(ST_Length(geom))/1000
[sdbms=# From nyc_streets;
    ?column?
-----
10418.904717199996
(1 row)
```

5. How long is ‘Columbus Cir’ (Columbus Circle)?

```
SELECT ST_Length(geom)
FROM nyc_streets
WHERE name = 'Columbus Cir';
```

```
[sdbms=# Select ST_Length(geom)
[sdbms=# From nyc_streets
[sdbms=# Where name='Columbus Cir';
    st_length
-----
308.3419936909855
(1 row)
```

6. How many polygons are in the ‘West Village’ multipolygon?

```
SELECT ST_NumGeometries(geom)
FROM nyc_neighborhoods
WHERE name = 'West Village';
```

```
[sdbms=# Select ST_NumGeometries(geom)
[sdbms=# From nyc_neighborhoods
[sdbms=# Where name='West Village';
  st_numgeometries
-----
      1
(1 row)
```

7. What is the length of streets in New York City, summarized by type?

```
SELECT      type, Sum(ST_Length(geom)) AS length
FROM        nyc_streets
GROUP BY    type
ORDER BY    length DESC;
```

```
[sdbms=# Select type,Sum(ST_Length(geom)) As length
[sdbms=# from nyc_streets
[sdbms=# group by type
[sdbms=# order by length DESC;
          type           |      length
-----+-----
residential      | 8629870.33786606
motorway         | 403622.4781263628
tertiary          | 360394.8790513027
motorway_link    | 294261.419479668
secondary         | 276264.3038979258
unclassified     | 166936.37160445828
primary           | 135034.2330179469
footway           | 71798.48783780965
service            | 28337.635038596007
trunk              | 20353.58198260764
cycleway           | 8863.751448259294
pedestrian         | 4867.050328250257
construction       | 4803.081621035617
residential; motorway_link | 3661.5750629374543
trunk_link         | 3202.1898124020076
primary_link       | 2492.5745708353616
living_street      | 1894.6390545733245
primary; residential; motorway_link; residential | 1367.7657694133486
undefined           | 380.5386191034604
steps               | 282.74522134212725
motorway_link; residential | 215.07778911517033
(21 rows)
```

Spatial Relationship Function

1. **ST_Equals(geometry A, geometry B)** tests the spatial equality of two geometries. It returns TRUE if two geometries of the same type have identical x,y coordinate values, i.e. if the second shape is equal (identical) to the first shape.

The ordering of points can be different but represent the same geometry structure

Example:

```
SELECT name, geom, ST_AsText(geom)
FROM   nyc_subway_stations
WHERE  name = 'Broad St';
```

```

sdbms=# Select name,geom,ST_AsText(geom)
sdbms=# From nyc_subway_stations
sdbms=# Where name='Broad St';
      name          |           geom           |           st_astext
-----+-----+-----+
 Broad St | 0101000020266900000EEBD4CF27CF2141BC17D69516315141 | POINT(583571.9059213118 4506714.341192182)
(1 row)

```

```

SELECT name
FROM nyc_subway_stations
WHERE ST_Equals(geom,
'0101000020266900000EEBD4CF27CF2141BC17D69516315141');

```

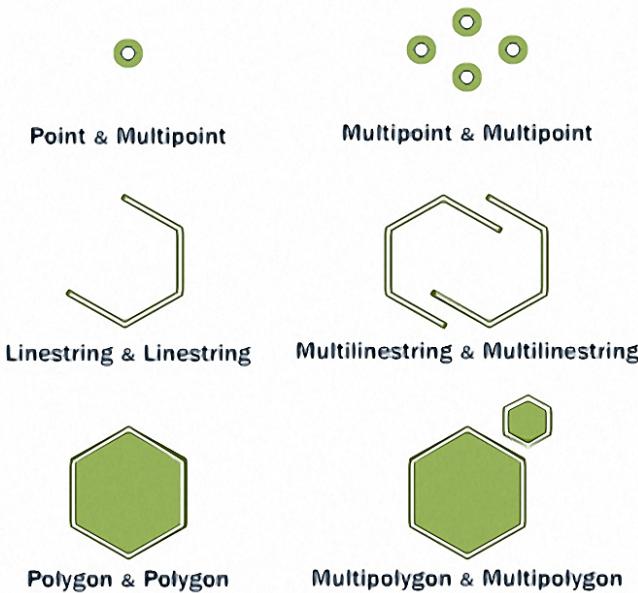
This code is an exact representation of a specific point. It is not very human readable. But for a test like equality, using the exact coordinates is necessary.

```

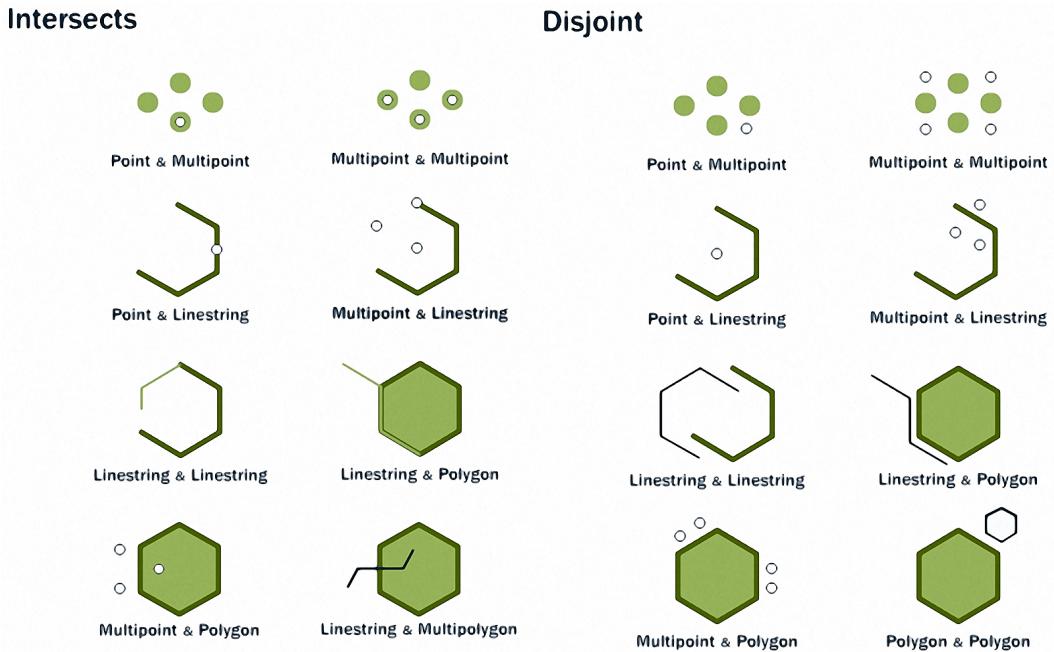
sdbms=# Select name
sdbms=# From nyc_subway_stations
sdbms=# Where ST_Equals(geom, '0101000020266900000EEBD4CF27CF2141BC17D69516315141');
      name
-----
 Broad St
(1 row)

```

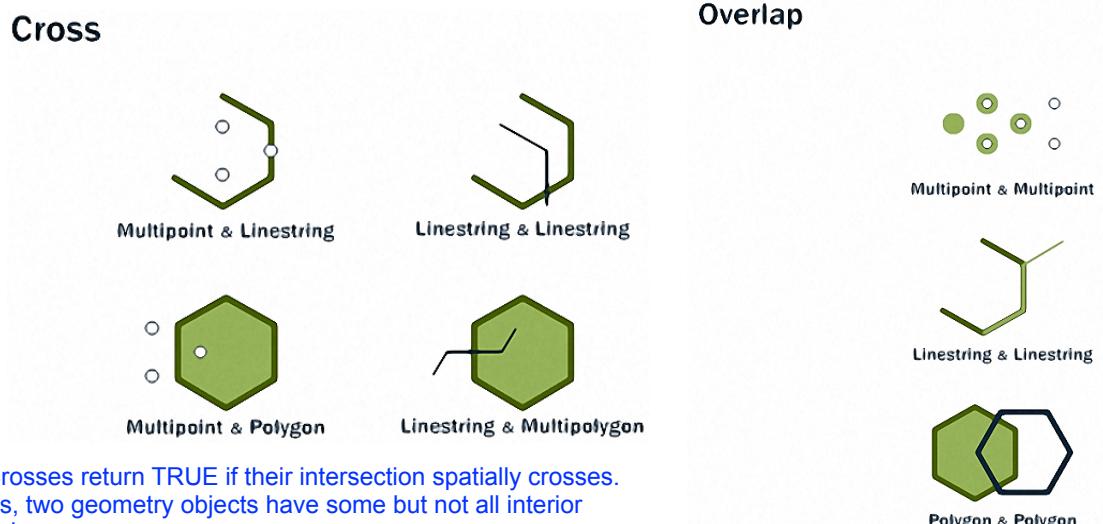
Equals



2. **ST_Intersects(geometry A, geometry B)** returns t (TRUE) if the two shapes have any space in common, i.e., if their boundaries or interiors intersect.
3. **ST_Disjoint(geometry A, geometry B)** is the Opposite of ST_INTERSECTS. If two geometries are disjoint, they do not intersect, and vice-versa. In fact, it is often more efficient to test “not intersects” than to test “disjoint” because the intersects tests can be spatially indexed, while the disjoint test cannot.



4. **ST_Crosses(geometry A, geometry B)** returns t (TRUE) if the intersection results in a geometry whose dimension is one less than the maximum dimension of the two source geometries and the intersection set is interior to both source geometries.
5. **ST_Overlaps(geometry A, geometry B)** compares two geometries of the same dimension and returns TRUE if their intersection set results in a geometry different from both but of the same dimension.



Example:

```
SELECT name, ST_AsText(geom)
FROM nyc_subway_stations
WHERE name = 'Broad St';
```

```

sdbms=# Select name,ST_AsText(geom)
sdbms=# From nyc_subway_stations
sdbms=# Where name='Broad St';
      name   |          st_astext
-----+-----
 Broad St | POINT(583571.9059213118 4506714.341192182)
(1 row)

```

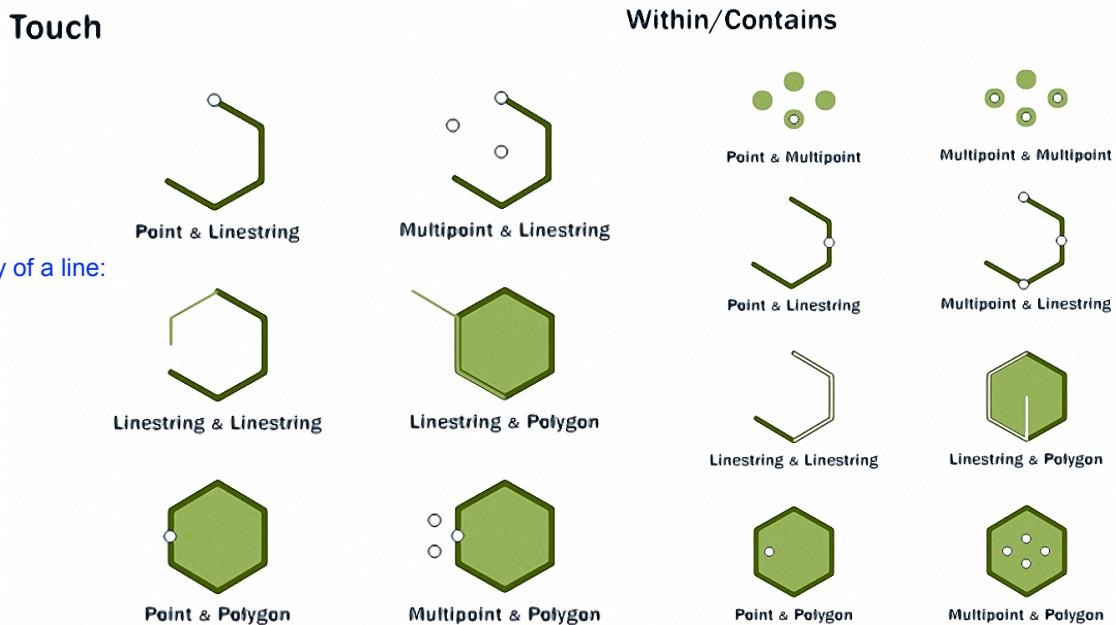
```

SELECT name, boroname
FROM nyc_neighborhoods
WHERE ST_Intersects(geom, ST_GeomFromText('POINT(5835714506714)', 26918));
sdbms=# Select name,boroname
sdbms=# From nyc_neighborhoods
sdbms=# Where ST_Intersects(geom,ST_GeomFromText('POINT(583571 4506714)',26918));
      name   |   boroname
-----+-----
 Financial District | Manhattan
(1 row)

```

This query returns the name and boroname of every neighborhood that intersects with the Broad St subway station.

6. **ST_Touches** tests whether two geometries touch at their boundaries, but do not intersect in their interiors. **ST_Touches(geometry A, geometry B)** returns TRUE if either of the geometries' boundaries intersect or if only one of the geometry's interiors intersects the other's boundary.
7. **ST_Within(geometry A , geometry B)** returns TRUE if the first geometry is completely within the second geometry. **ST_Within** tests for the exact opposite result of **ST_Contains**.
8. **ST_Contains(geometry A, geometry B)** returns TRUE if the second geometry is completely contained by the first geometry.

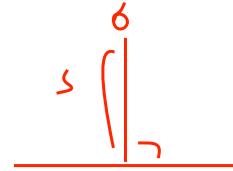


9. The **ST_Distance(geometry A, geometry B)** calculates the shortest distance between two geometries and returns it as a float. This is useful for actually reporting back the distance between objects.

Example:

```
SELECT ST_Distance(  
    ST_GeometryFromText('POINT(0 5)'),  
    ST_GeometryFromText('LINESTRING(-2 2, 2 2)'));
```

```
sdbms=# Select ST_Distance(  
sdbms# ST_GeometryFromText('POINT(0 5)'),  
sdbms# ST_GeometryFromText('LINESTRING(-2 2,2 2)'));  
st_distance  
-----  
3  
(1 row)
```



10. For testing whether two objects are within a distance of one another, the **ST_DWithin** function provides an index-accelerated true/false test. This is useful for questions like “how many trees are within a 500 meter buffer of the road?”. You don’t have to calculate an actual buffer, you just have to test the distance relationship.

ST_Dwithin

ST_DWithin(geometry A, geometry B, radius):
returns TRUE if geometry A is radius distance or less from geometry B



Point & Point (True)



Point & Point (False)



Polygon & Point (True)



Polygon & Point (False)

Example:

```
SELECT name  
FROM nyc_streets  
WHERE ST_DWithin(geom, ST_GeomFromText('POINT(583571 4506714)',  
26918), 10);
```

```
sdbms=# Select name  
sdbms# From nyc_streets  
sdbms# Where ST_DWithin(geom,ST_GeomFromText('POINT(583571 4506714)', 26918), 10);  
name  
-----  
Wall St  
Broad St  
Nassau St  
(3 rows)
```

Spatial Relationship Query

1. “What is the geometry value for the street named ‘Atlantic Commons’?”

```
SELECT ST_AsText(geom)
FROM nyc_streets
WHERE name = 'Atlantic Commons';
```

```
sdbms=# SELECT ST_AsText(geom)
sdbms=# From nyc_streets
sdbms=# Where name='Atlantic Commons';
                           st_astext
-----
 MULTILINESTRING((586781.7015777241 4504202.153143394,586863.5196448397 4504215.988170098))
(1 row)
```

2. “What neighborhood and borough is Atlantic Commons in?”

```
SELECT name, boroname
FROM nyc_neighborhoods
WHERE ST_Intersects(geom, ST_GeomFromText('LINESTRING(586782
4504202,586864 4504216)', 26918));
```

[Is ST_Touches\(\) okay for this query?](#)

```
sdbms=# Select name,boroname
sdbms=# From nyc_neighborhoods
sdbms=# Where ST_Intersects(geom,ST_GeomFromText('LINESTRING(586782 4504202,586864 4504216)', 26918));
      name   | boroname
-----
 Fort Green | Brooklyn
(1 row)
```

3. “What streets does Atlantic Commons join with?”

```
SELECT name
FROM nyc_streets
WHERE ST_DWithin(geom,ST_GeomFromText('LINESTRING(586782
4504202,586864 4504216)', 26918), 0.1);
```

[distance threshold = 0.1 ?](#)
Again, the coordinates of the linestring have been rounded rather than the exact values. So it is not proper to use ST_Crosses() here.

```
sdbms=# Select name
sdbms=# From nyc_streets
sdbms=# Where ST_DWithin(geom,ST_GeomFromText('LINESTRING(586782 4504202,586864 4504216)', 26918), 0.1);
      name
-----
 Atlantic Commons
 Cumberland St
(2 rows)
```

4. “Approximately how many people live on (within 50 meters of) Atlantic Commons?”

```
SELECT Sum(popn_total)
FROM nyc_census_blocks
WHERE ST_DWithin(geom, ST_GeomFromText('LINESTRING(586782
4504202,586864 4504216)', 26918),50);
```

```
sdbms=# Select Sum(popn_total)
sdbms=# From nyc_census_blocks
sdbms=# Where ST_DWithin(geom,ST_GeomFromText('LINESTRING(586782 4504202,586864 4504216)', 26918),50);
      sum
-----
 1438
(1 row)
```

Spatial Join Query

By default, we are using an INNER JOIN.

In the previous section, we explored spatial relationships using a two-step process: first we extracted a subway station point for ‘Broad St’; then, we used that point to ask further questions such as “what neighborhood is the ‘Broad St’ station in?”

Using a spatial join, we can answer the question in one step, retrieving information about the subway station and the neighborhood that contains it:

[Any function that provides a true/false relationship between two tables can be used to drive a spatial join, but the most commonly used ones are: ST_Intersects, ST_Contains, and ST_DWithin.](#)

```
SELECT      subways.name AS subway_name,
            neighborhoods.name AS neighborhood_name,
            neighborhoods.borongame AS borough
FROM        nyc_neighborhoods AS neighborhoods
JOIN        nyc_subway_stations AS subways
ON          ST_Contains(neighborhoods.geom, subways.geom)
WHERE       subways.name = 'Broad St';
```

```
sdbms=# Select subways.name As subway_name,
sdbms-# neighborhoods.name As neighborhood_name,
sdbms-# neighborhoods.borongame As borough
sdbms-# From nyc_neighborhoods As neighborhoods
sdbms-# Join nyc_subway_stations As subways
sdbms-# On ST_Contains(neighborhoods.geom, subways.geom)
sdbms-# Where subways.name='Broad St';
subway_name | neighborhood_name | borough
-----+-----+-----+
Broad St    | Financial District | Manhattan
(1 row)
```

1. “What is the population and racial make-up of the neighborhoods of Manhattan?”

```
SELECT      neighborhoods.name AS neighborhood_name,
            Sum(census.popn_total) AS population,
            100.0 * Sum(census.popn_white)/Sum(census.popn_total) AS white_pct,
            100.0 * Sum(census.popn_black) / Sum(census.popn_total) AS black_pct
FROM        nyc_neighborhoods AS neighborhoods
JOIN        nyc_census_blocks AS census
ON          ST_Intersects(neighborhoods.geom, census.geom)
WHERE       neighborhoods.borongame = 'Manhattan'
GROUP BY   neighborhoods.name
ORDER BY   white_pct DESC;
```

1. The JOIN clause creates a virtual table that includes columns from both the neighborhoods and census tables.
2. The WHERE clause filters our virtual table to just rows in Manhattan.
3. The ON clause helps to remain all the geometry areas that interact between two tables
4. Sum up the population and format (e.g., GROUP BY, ORDER BY) on the final numbers

```

sdbms=# SELECT neighborhoods.name AS neighborhood_name,
sdbms#           Sum(census.popn_total) AS population,
sdbms#           100.0 * Sum(census.popn_white)/Sum(census.popn_total) AS white_pct,
sdbms#           100.0 * Sum(census.popn_black) / Sum(census.popn_total) AS black_pct
sdbms# FROM      nyc_neighborhoods AS neighborhoods
sdbms# JOIN      nyc_census_blocks AS census
sdbms# ON        ST_Intersects(neighborhoods.geom, census.geom)
sdbms# WHERE     neighborhoods.borname = 'Manhattan'
sdbms# GROUP BY  neighborhoods.name
sdbms# ORDER BY   white_pct DESC;

```

neighborhood_name	population	white_pct	black_pct
Carnegie Hill	18763	90.08154346319886	1.4123541011565315
West Village	26718	87.59637697432443	2.1745639643685903
North Sutton Area	22460	87.56455921638468	1.553873552983081
Upper East Side	203741	85.02216048807064	2.7220834294520984
Soho	15436	84.64628142005701	2.235034983156258
Greenwich Village	57224	81.97609394659584	2.43778834055641
Central Park	46600	79.49356223175965	7.967811158798283
Tribeca	20908	79.1180409412665	3.548880811172757
Gramercy	104876	75.4586368663946	4.720813150768526
Murray Hill	29655	75.01264542235711	2.5122239082785365
Chelsea	61340	74.81578089338116	6.372676882947506
Upper West Side	214761	74.561489283436	9.188353565125885
Midtown	76840	72.60150963040083	5.166579906298803
Battery Park	17153	71.8300005829884	3.375502827493733
Financial District	34807	69.92846266555578	3.829689430287011
Clinton	32201	65.2743703611689	7.943852675382752
East Village	82266	63.26550458269516	8.771545960664186
Garment District	10539	55.18550147072777	7.08795900939368
Morningside Heights	42844	52.719167211278126	19.370273550555503
Little Italy	12568	49.029280712921704	1.8220878421387652
Yorkville	58450	35.5551753635586	29.72284003421728
Inwood	50047	35.21489799588387	16.846164605271046
Washington Heights	169013	34.873057102116405	16.827107973942834
Lower East Side	96156	33.50988939639752	9.075876700361912
East Harlem	60576	26.43621236133122	40.38398045430534
Hamilton Heights	67432	23.889251393996915	35.772333610155414
Chinatown	16209	15.16441483126658	3.757171941513974
Harlem	134955	15.094661183357415	67.06161313030269
(28 rows)			

2. “What subway station is in ‘Little Italy’? What subway route is it on?”

```

SELECT      s.name, s.routes
FROM        nyc_subway_stations AS s
JOIN        nyc_neighborhoods AS n
ON          ST_Contains(n.geom, s.geom)      Find the stations contained by
WHERE       n.name = 'Little Italy';          the neighborhood called 'Little Italy'

```

```

sdbms=# SELECT s.name, s.routes
sdbms# FROM    nyc_subway_stations AS s
sdbms# JOIN    nyc_neighborhoods AS n
sdbms# ON      ST_Contains(n.geom, s.geom)
sdbms# WHERE   n.name = 'Little Italy';

```

name	routes
Spring St	6

(1 row)

3. “What are all the neighborhoods served by the 6-train?”

```
SELECT      DISTINCT n.name, n.borongame
FROM        nyc_subway_stations AS s
JOIN        nyc_neighborhoods AS n
ON          ST_Contains(n.geom, s.geom)
WHERE       strpos(s.routes,'6') > 0;
```

the DISTINCT keyword is to remove duplicate values from result set where there were more than one subway station in a neighborhood.

strpos(str, substr): find the position from where the substring is being matched within the string

```
sdbms=# SELECT  DISTINCT n.name, n.borongame
sdbms-# FROM    nyc_subway_stations AS s
sdbms-# JOIN    nyc_neighborhoods AS n
sdbms-# ON      ST_Contains(n.geom, s.geom)
sdbms-# WHERE   strpos(s.routes,'6')>0;
      name      | borongame
-----+-----
Chinatown      | Manhattan
East Harlem    | Manhattan
Financial District | Manhattan
Gramercy       | Manhattan
Greenwich Village | Manhattan
Hunts Point    | The Bronx
Little Italy   | Manhattan
Midtown        | Manhattan
Mott Haven     | The Bronx
Murray Hill    | Manhattan
Parkchester    | The Bronx
Soundview      | The Bronx
South Bronx    | The Bronx
Upper East Side | Manhattan
Yorkville      | Manhattan
(15 rows)
```

4. “After 9/11, the ‘Battery Park’ neighborhood was off limits for several days. How many people had to be evacuated?”

```
SELECT      Sum(popn_total)
FROM        nyc_neighborhoods AS n
JOIN        nyc_census_blocks AS c
ON          ST_Intersects(n.geom, c.geom)
WHERE       n.name = 'Battery Park';
```

```
sdbms=# SELECT  Sum(popn_total)
sdbms-# FROM    nyc_neighborhoods AS n
sdbms-# JOIN    nyc_census_blocks AS c
sdbms-# ON      ST_Intersects(n.geom, c.geom)
sdbms-# WHERE   n.name = 'Battery Park';
      sum
-----
  17153
(1 row)
```

5. “What are the population density (people / km²) of the ‘Upper West Side’ and ‘Upper East Side’?”

```
SELECT      n.name,
              Sum(c.popn_total) / (ST_Area(n.geom) / 1000000.0) AS popn_per_sqkm
FROM        nyc_census_blocks AS c
JOIN         nyc_neighborhoods AS n
ON          ST_Intersects(c.geom, n.geom)
WHERE       n.name = 'Upper West Side'
OR           n.name = 'Upper East Side'
GROUP BY    n.name, n.geom;
```

name	popn_per_sqkm
Upper East Side	48524.48774898572
Upper West Side	40152.48960800237

(2 rows)