



Q & A session for INFS 4205/7205

-- A brief review

- Fengmei Jin
- fengmei.jin@uq.edu.au

+ Spatial Data



- Any data with a location component
 - 2D Space
 - Geographical space: GIS, Urban Planning
 - 3D Space
 - (x, y, z) : The universe, brain model, molecule structure...
 - (x, y, t) : Trajectory
- Two types of spatial data
 - Those data about the space (e.g., road networks, maps)
 - Those data about objects (e.g., location of shops, location about cars)

+ Spatial DBMS

- A spatial database system is a **database system**
- It offers **spatial data types** in its data model and the **query language**
- It supports spatial data types in its implementation, by providing **spatial indexing** and **efficient algorithms** for **spatial join queries** and other types of spatial queries
 - Retrieve the spatial data without scanning the whole set

+ Spatial Data Model in Oracle

- Element, Geometry and Layer
- **Element:** the basic building block of a geometric feature
 - **Point data:** One point, stored as an (x, y) pair
 - **Line data:** Two points representing the start and the end of a line segment
 - **Polygon data:** A sequence of coordinates, one vertex pair for each line segment of the polygon
 - Simple Polygon: Both boundary and the interior
 - Complex polygons (Geometry)
 - Self-intersecting boundary
 - Multiple dis-connected components




+ Spatial Data Model in Oracle

- **Geometry:** representation of a user's spatial feature, modelled as an ordered **set of elements**
 - Each geometry has a unique ID, and can be associated with a set of attributes
 - A geometry might describe a lake
 - A polygon with nested polygons for islands
 - Attributes such as lake name, water capacity, fauna, flora, ...



+ Spatial Data Model in Oracle

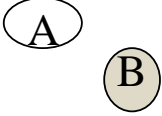
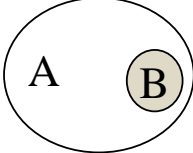
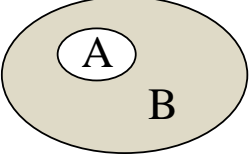
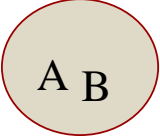

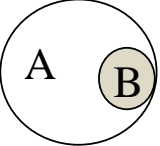
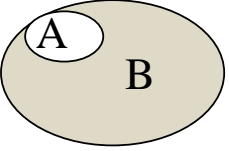
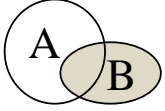
- **Layer:** a collection of geometries having the same attribute set
- Examples: soil types, road network, political boundaries, population density, crops, weather conditions,...



+ Spatial Relationships

■ Topological

- Invariant under translation such as rotation and scaling

 <p>disjoint</p>	 <p>contain</p>	 <p>inside</p>	 <p>equal</p>
 <p>meet</p>	 <p>cover</p>	 <p>covered_by</p>	 <p>overlap</p>


■ Directional

- E.g., Above, Left, North of,...
- May change with rotation

■ Metric

- E.g., Distance, Length, Area,...
- May change with scaling

+ The 9-Intersection Matrix



A_b : boundary

A_i : interior

A_e : exterior

$$\begin{pmatrix} A_b \cap B_b & A_b \cap B_i & A_b \cap B_e \\ A_i \cap B_b & A_i \cap B_i & A_i \cap B_e \\ A_e \cap B_b & A_e \cap B_i & A_e \cap B_e \end{pmatrix}$$

Each element is either 1 or 0.

+ Spatial Indexing

- Purpose:

- Efficiency in processing spatial selection, join and other spatial operations

- Two strategies to organize space and objects


- Map spatial objects into 1D space and use a standard index structure (B-tree)
- Dedicated external data structures

- Basic ideas

- Approximation
 - Bounding box, Grids
- Hierarchical Data Organization

+ Object Approximation

- A fundamental idea of spatial indexing is the use of approximation
- Continuous Approximation
 - Object centric
 - Example:
 - Use of MBRs (Minimum Bounding Rectangles)
 - R-Tree
- Grid Approximation
 - Space centric
 - Faster mapping
 - Uniform / Non-uniform
 - High-D?
 - Example:
 - Quad-Tree



+ Line Data

- Arbitrary direction and shape
 - Use MBR and treat like polygons
 - Treated as trajectory data
- Straight line segments with 'perpendicular directions'
 - Align parallel the axis
 - Interval Tree
- Straight line segments with different directions
 - Use Point Index techniques to index the two end points separately
 - Segment Tree

+ Indexing Structures for line/segment query

■ 2D Range Tree – Binary Tree of Binary Tree

- One Binary Tree in X (x-Tree)
- One Binary Tree in Y for each node in the X-Tree

■ Interval Tree

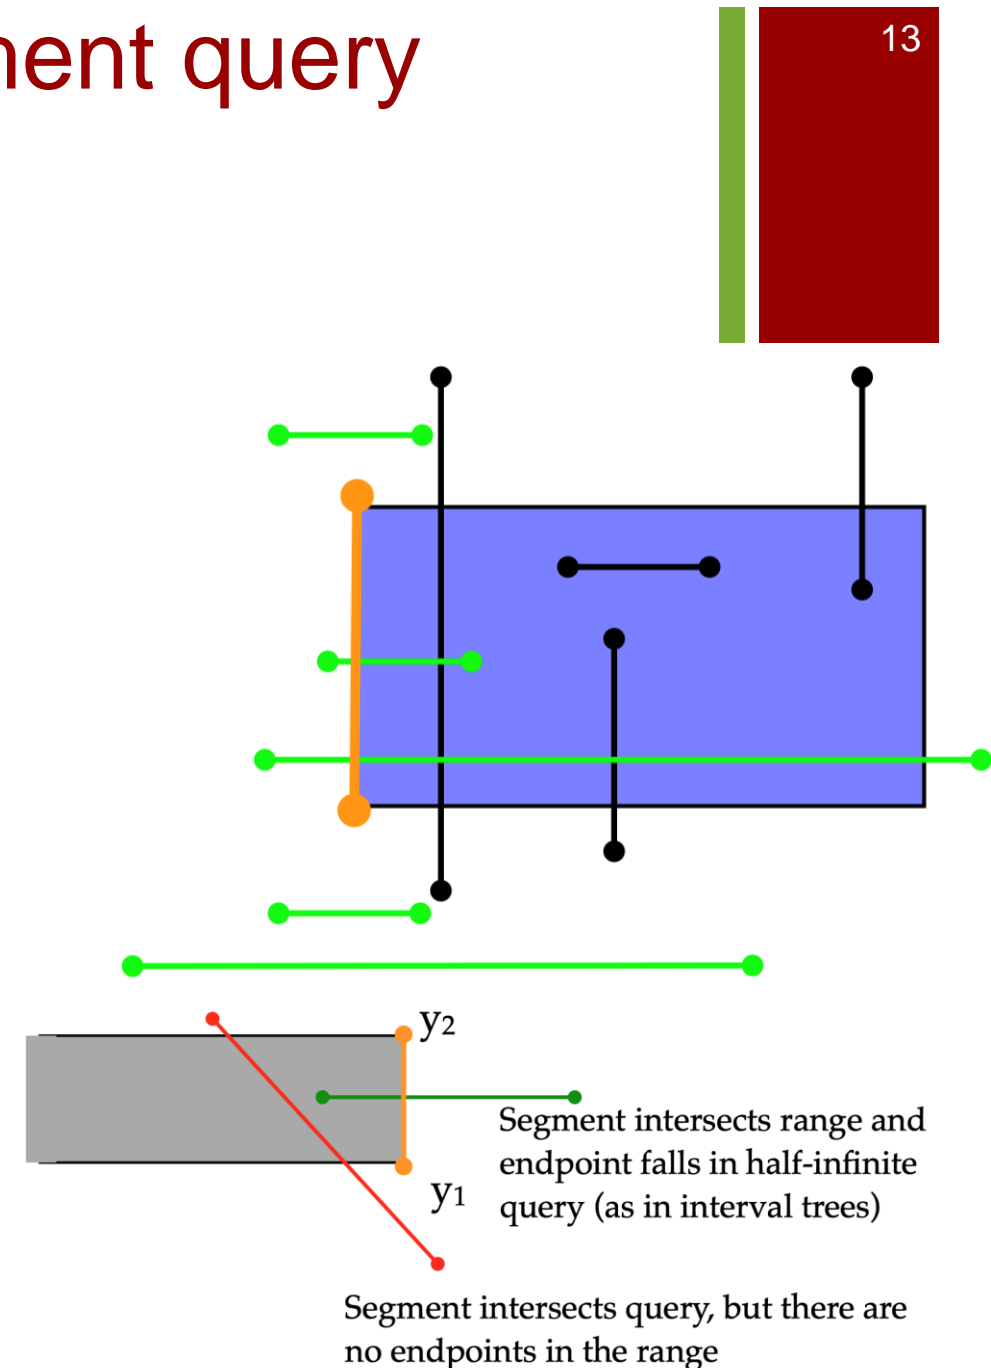
- Each node stores endpoints of intervals located inside
 - Two sorted lists / 2D range tree

■ Priority Search Tree

- Good for 1-Side Range Query: $[-\infty, x], [y_1, y_2]$
 - Heap for one dimension (like x-axis)
 - Binary Search Tree for the other (y-axis)

■ Segment Tree

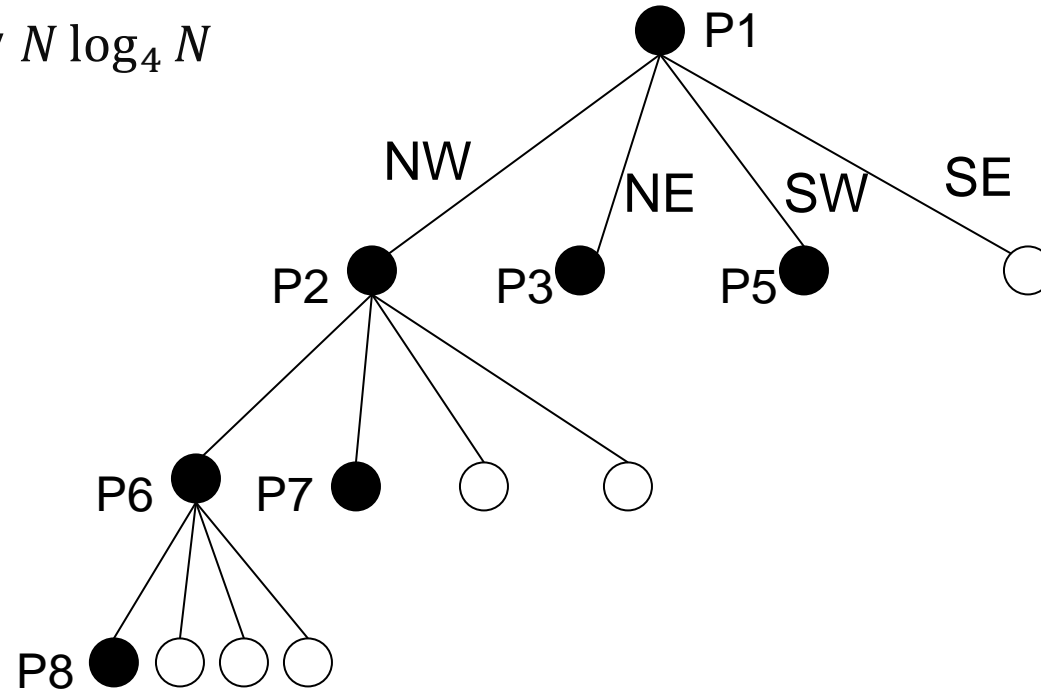

- Arbitrarily Oriented Segments
- Store the segments in the Binary Search Tree



+ Point Quadtree


■ Insertion

- Random insertion roughly $N \log_4 N$



■ When is the worst case?



- Insertion takes $N(N - 1)/2$



+ Region Quadtree



15

■ PR Quadtree



- Based on regular decomposition of the universe
 - Recursively decomposing a region into four congruent blocks
 - Only leaves contain data

+ kd-Tree Construction




Depends on the order of insertion (not robust for sorted data).

Variations: non-alternative, data at leaves only, representing regions etc.

+ Z-Order

■ How to obtain the z-order?

1. Counting: A is 24
 2. Quaternary: $(120)_4 = (24)_{10}$
 3. Bit-Interleaving
 - $x_0y_0x_1y_1\dots$
 - $(011000)_2 = (24)_{10}$
 - Works fine with varying resolutions
-
- B: $(21)_4$
 - $(1001)_2$



+ Z-Value Example

■ B Covers C (Base 4)


■ B: 21

■ C: 210

■ B covers C (Base 5)

■ B: 320

■ C: 321




+ Transformation: Using Z-Ordering

■ Granularity

- {11}, or {111, 112, 114}, or {111, 1121, 1123, 1124, 1141, 1142}

■ When decomposition stops

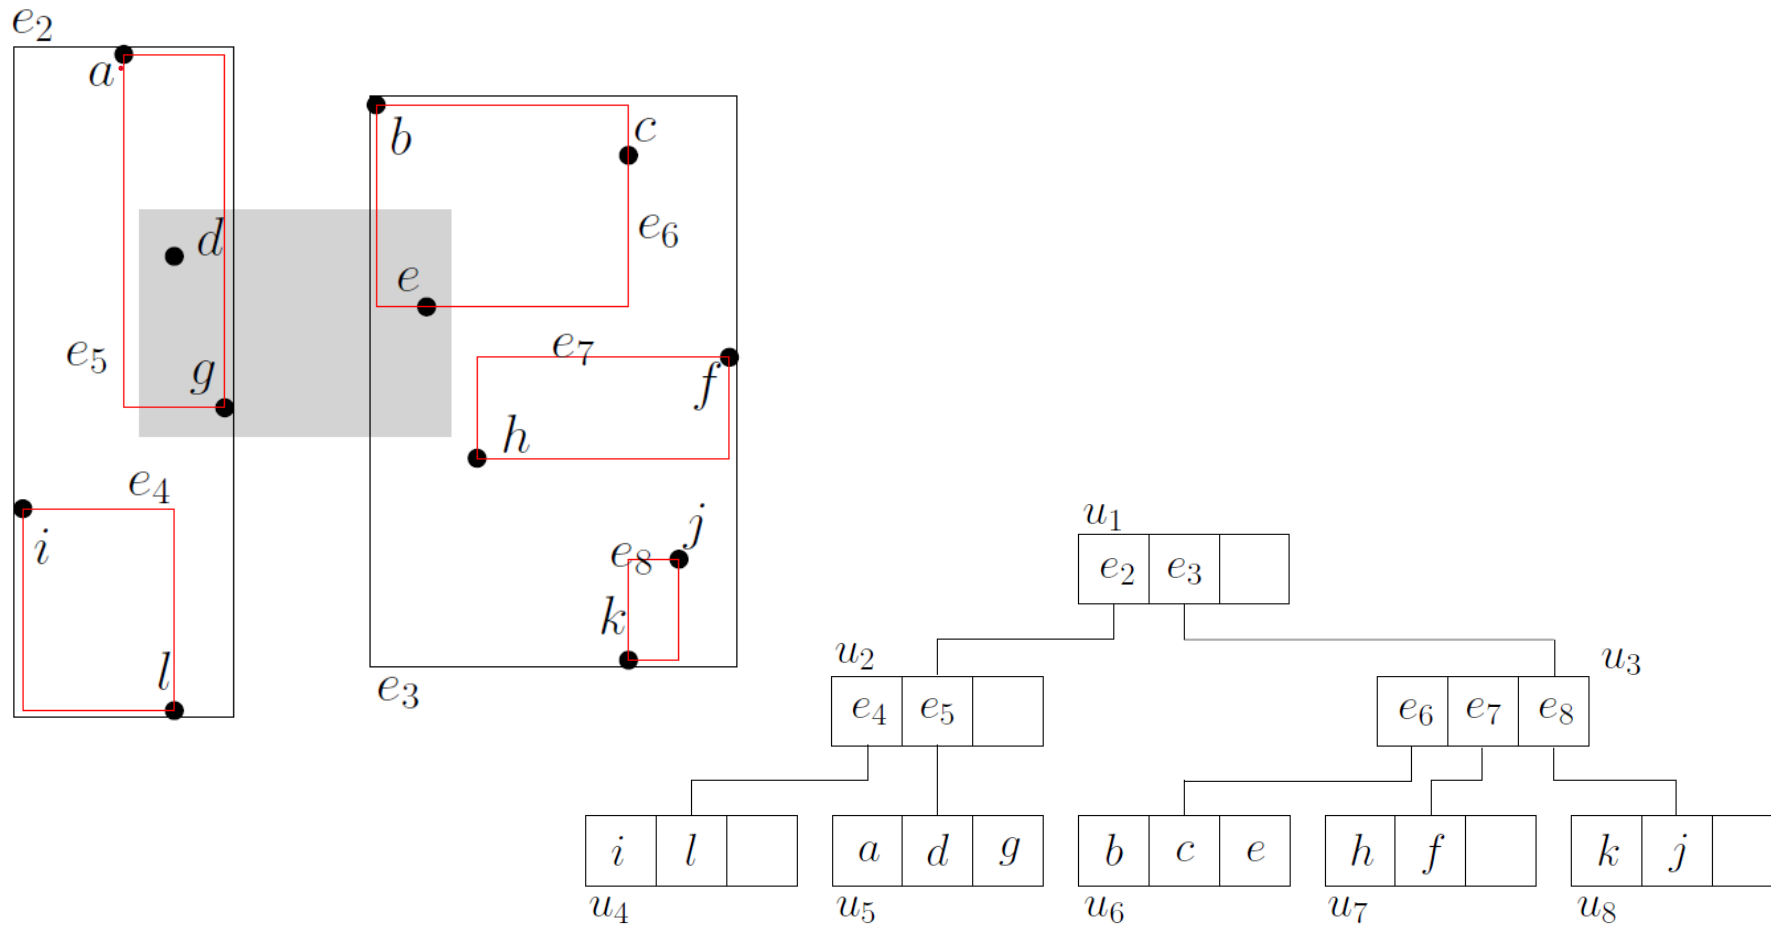
- Current cell either fully out or in the polygon
- Reached the “resolution”



...the entire space is 1.

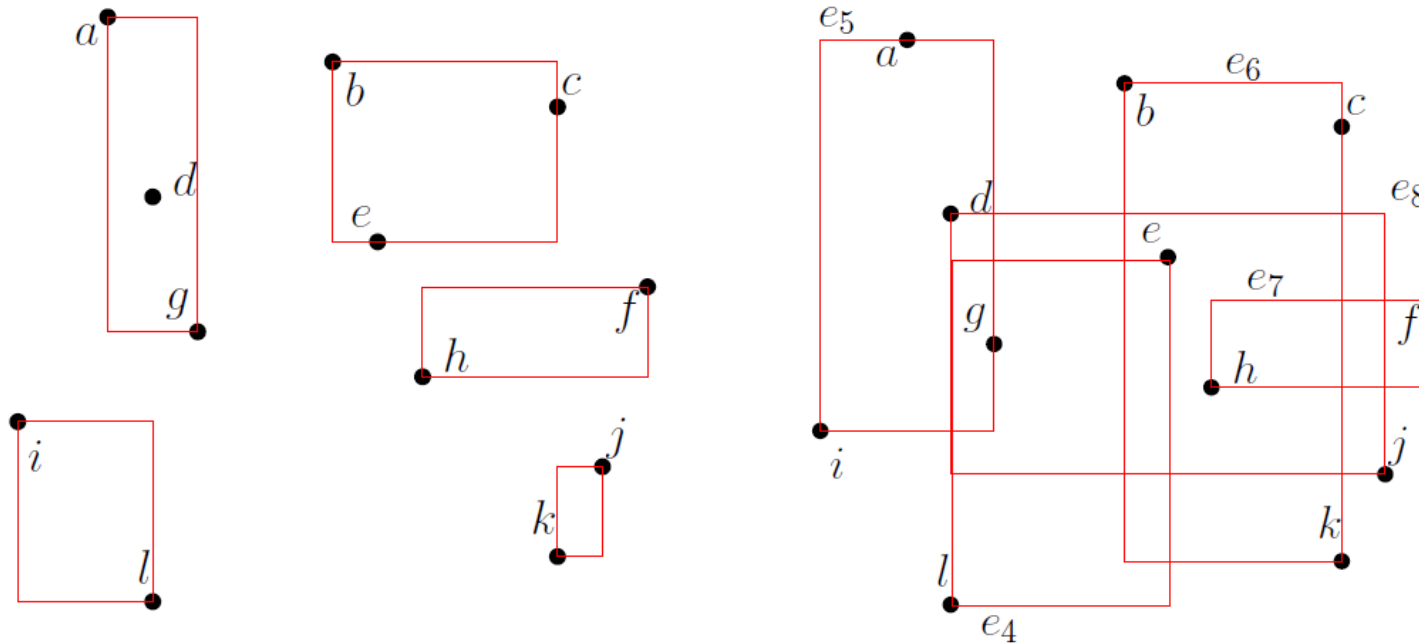
+ R-Tree Range Query

- u_1, u_2, u_3, u_5, u_6 are accessed



+ R-Tree Construction

- R-Tree construction can be “arbitrary”
 - Bottom-up
 - No formal constraint on the grouping of data into nodes





- The left tree has a smaller perimeter sum than the right one

+ Clipping

■ Motivation


- R-Tree: May examine all the MBRs at all levels
 - Because the MBR may overlap, the space is not disjointly decomposed
 - Query point Q in the example
- Single search path for a point query



+ R⁺-Tree

■ Basic ideas

- A hierarchy of overlapping MBRs → A hierarchy of disjoint MBRs
 - Regular grid / Irregular grid
- Clipping polygon at cell boundaries
 - Whenever an MBR at a lower level overlaps with another MBR, decompose it into a collection of non-overlapping sub-MBRs
- Allowing one polygon in multiple cells
 - Non-overlapping is achieved at the cost of space




+ Filter-and-Refine

- A most commonly used processing strategy
- Motivation
 - Avoid expensive spatial processing as much as possible
- Basic Idea
 - A filter step, followed by a refinement step
 - **Filter** step: applying *simple* operations on approximations of spatial objects
 - **Refinement** step: applying the *actual* spatial operations on the *full geometry* of spatial objects

+ Spatial Join Example

- Intersection join



- Join results: (x_2, y_1) , (x_2, y_2)

- Other spatial join operations

- Topological: intersection, adjacent, contains...
- Metrical or directional: within_distance...
- More advanced: nearest....

+ Processing Framework

■ Filter step

- Find a set of candidates $C = \{(p, s): p \in R \text{ and } s \in S\}$ *quickly*
 - Using approximations (e.g., MBR) and indexes
 - Other filter steps possible (eg, using *progressive* approximation)

■ Housekeeping step

- Process C such that the IO cost for the refinement step can be further minimized
- E.g., Removing duplicates; Performing refinement in optimal order

■ Refine step

- Fetch full geometry for the objects in each candidate, and apply a full test to drop “false hits”

+ Simple Nested-Loops Join

```
for each  $r$  in  $R$ 
  for each  $s$  in  $S$ 
    if ( $r, MBR$  intersect  $s.MBR$ )
      put ( $r, s$ ) to the candidate set;
```

+ Indexed Nested-Loops Join


- Using a window query against S

for each r in R

Find all s in S such that $s.MBR$ **intersect**
 $r.MBR$

put (r, s) to the candidate set;



+ Nested-Loops With R/R+-Trees



Is it possible to produce the same candidate (p_1, p_2) multiple times?


Reading: the paper by Brinkhoff, Kriegel and Seeger 1993.

+ Sort-Merge Join using Z-values



■ Differences

- 1. $131 \neq 130$
- 2. Cannot move from 130 to 223 immediately



Algorithm Sketch:

- 1) Two sorted lists and two pointers
- 2) Synchronized traversal
 - $\text{overlap}(r, s)$?
 - increase $\min(r, s)$
- 3) Some values in stack

Reading: the paper by Orenstein and Manola 1988.